



CLOUD COMPUTING WITH BIG DATA

By Group 8

Group members

秦尧
林志祥
吕哲坤
刘艺奇

Group Name: SkyNet Innovator

Pa 2024



Project Introduction

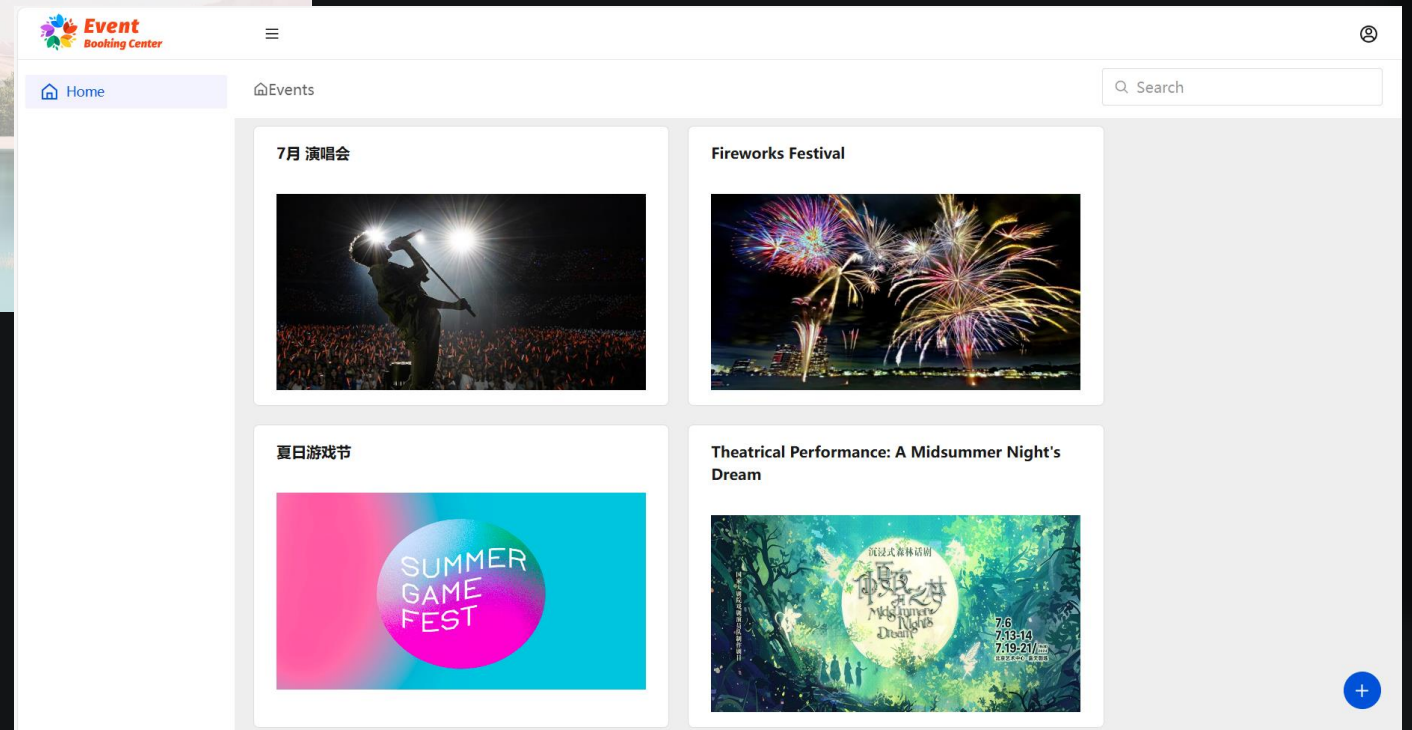
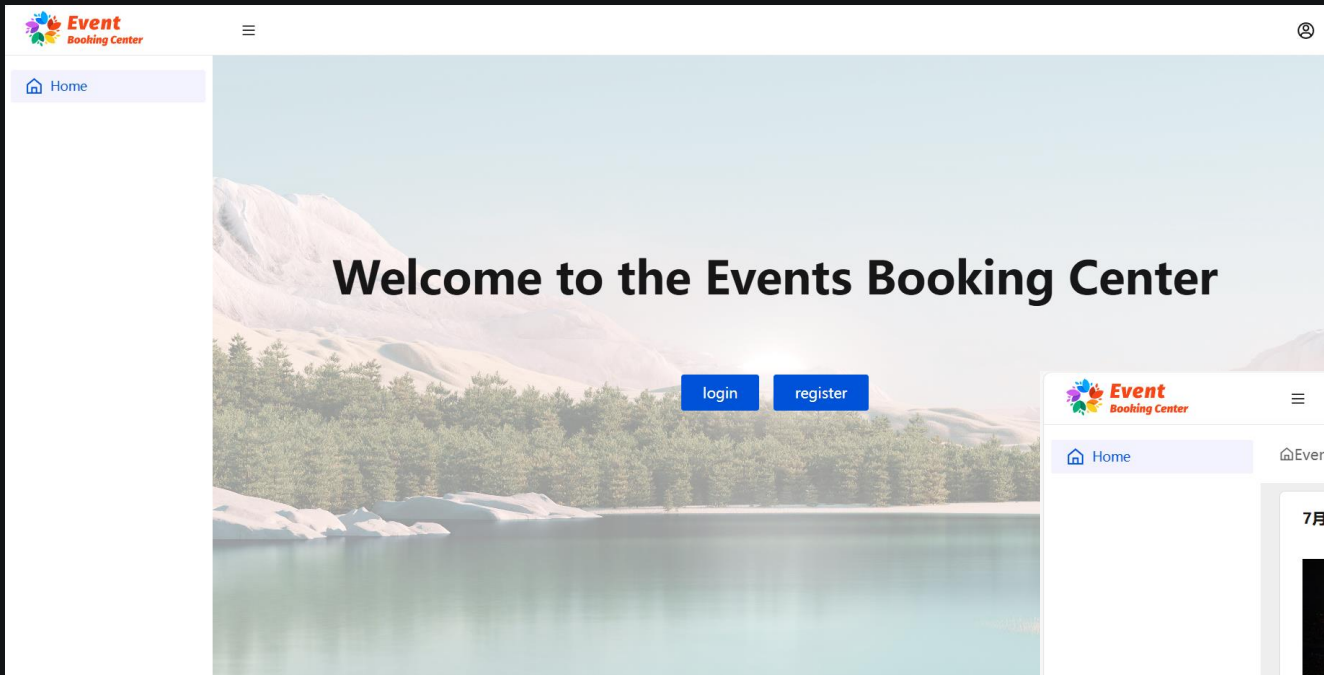


Event Booking Center

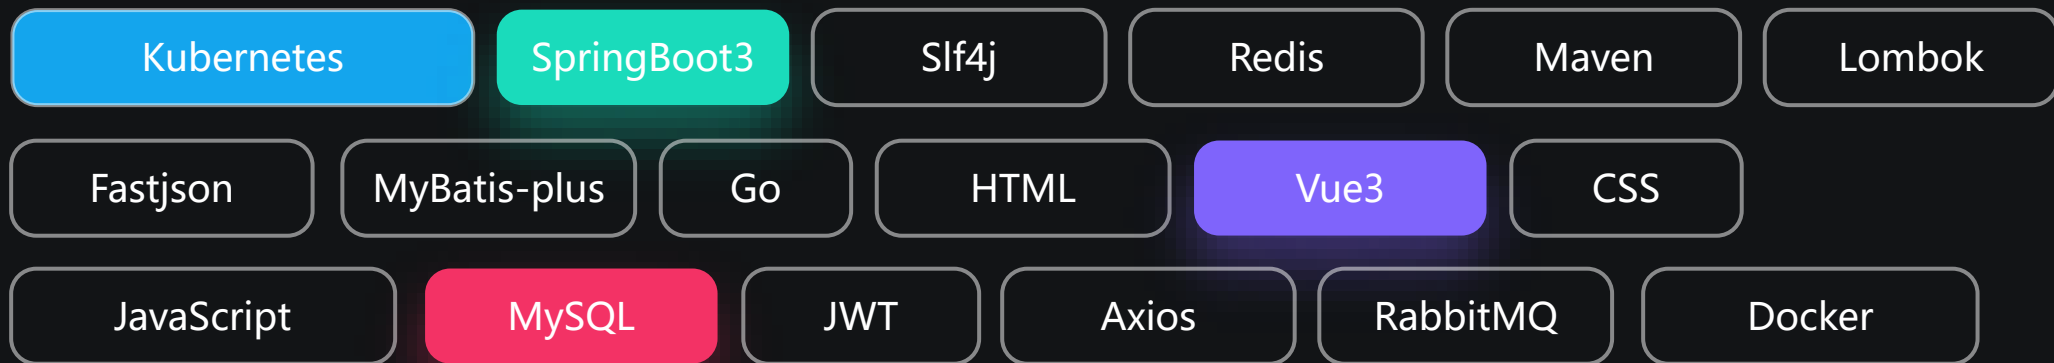
Everything
convenient

- Our event publishing and booking system allows users to easily manage activities and enables seamless scheduling of events
- Users can browse available slots, select their preferred time, and confirm their booking with just a few clicks.

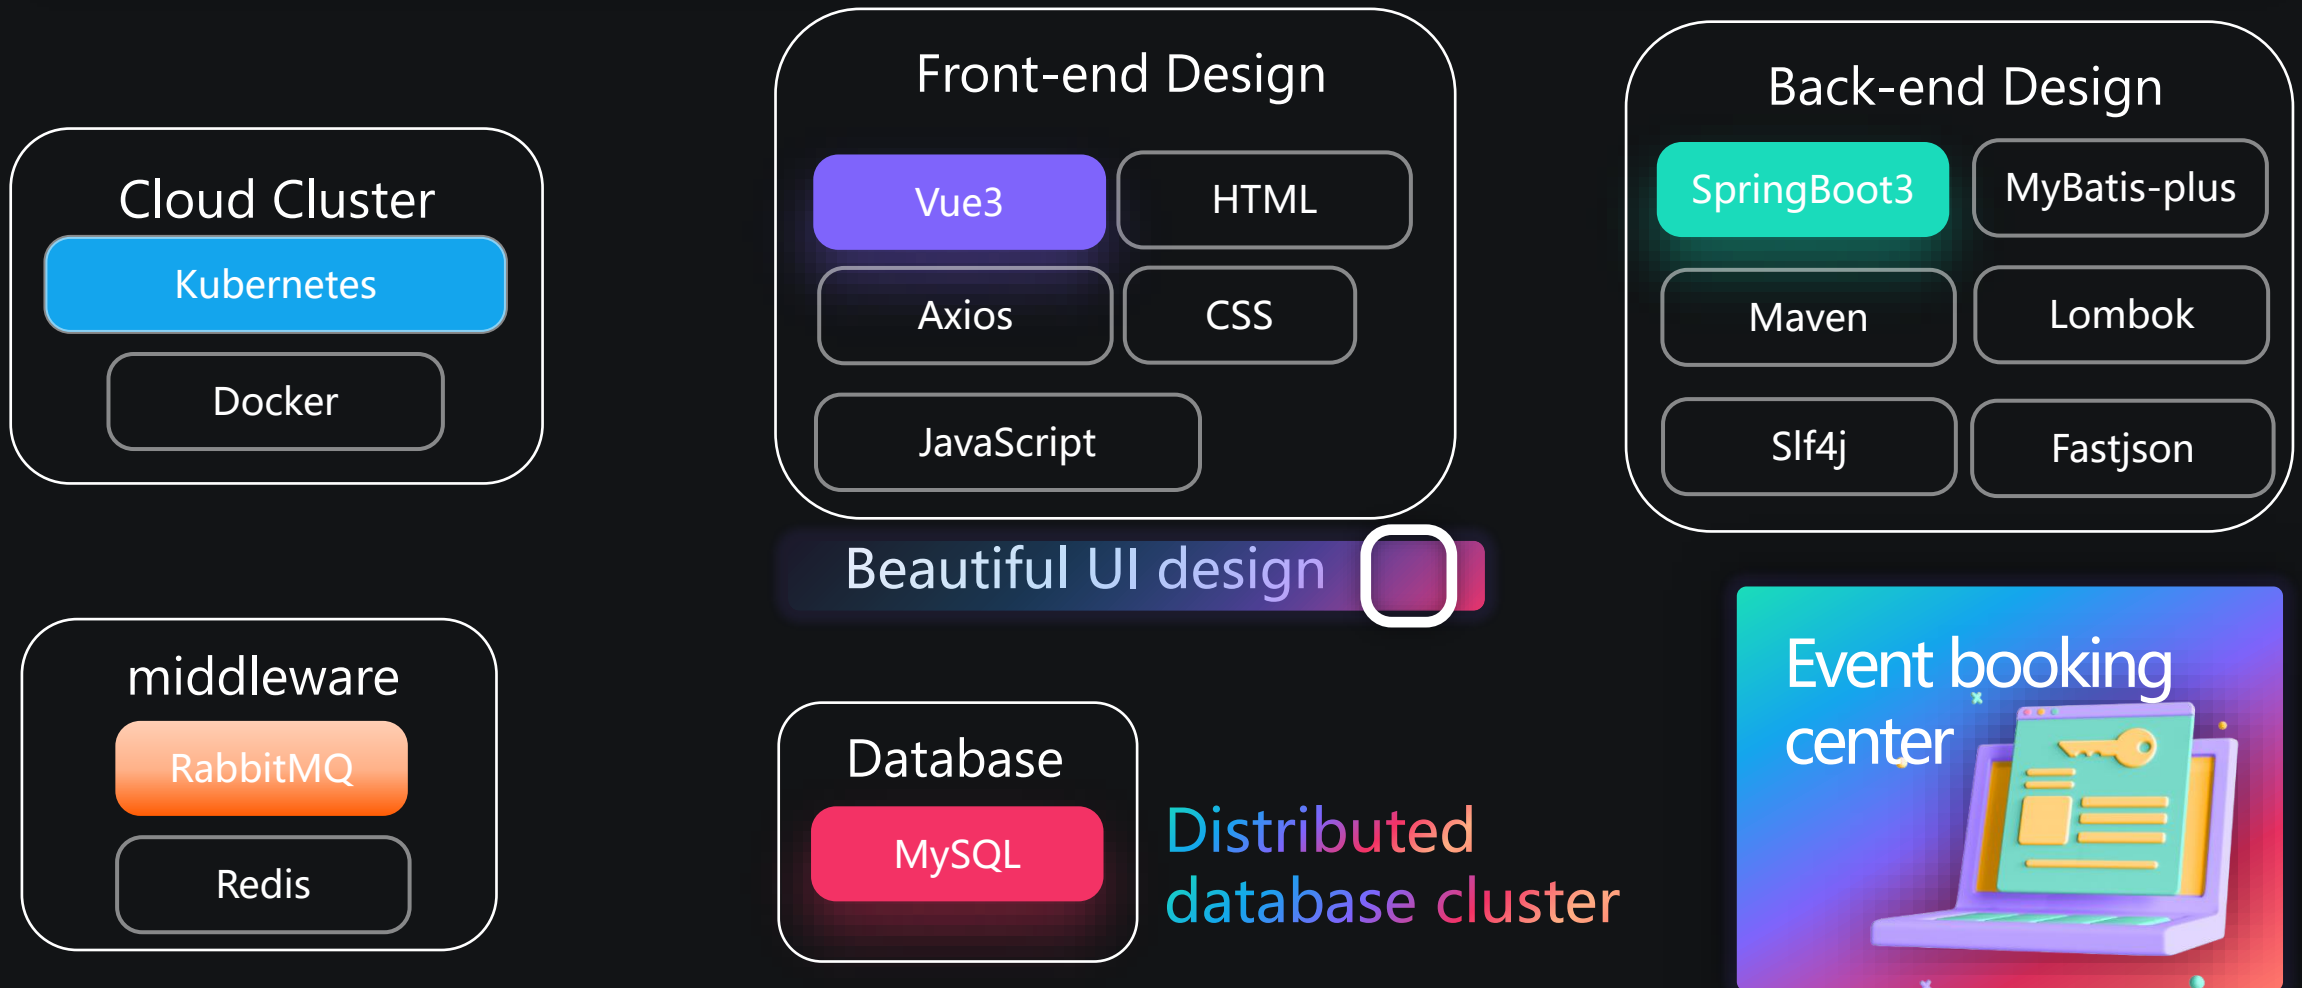
Front end display



Main technology stacks



Main technology stacks



Microservices Architecture

There are **9** Microservice including:

User Verification

userRoutingHashGenerateService

loginService

registerService

tokenVerificationService

Event Operation

orderRecordService

bookingService

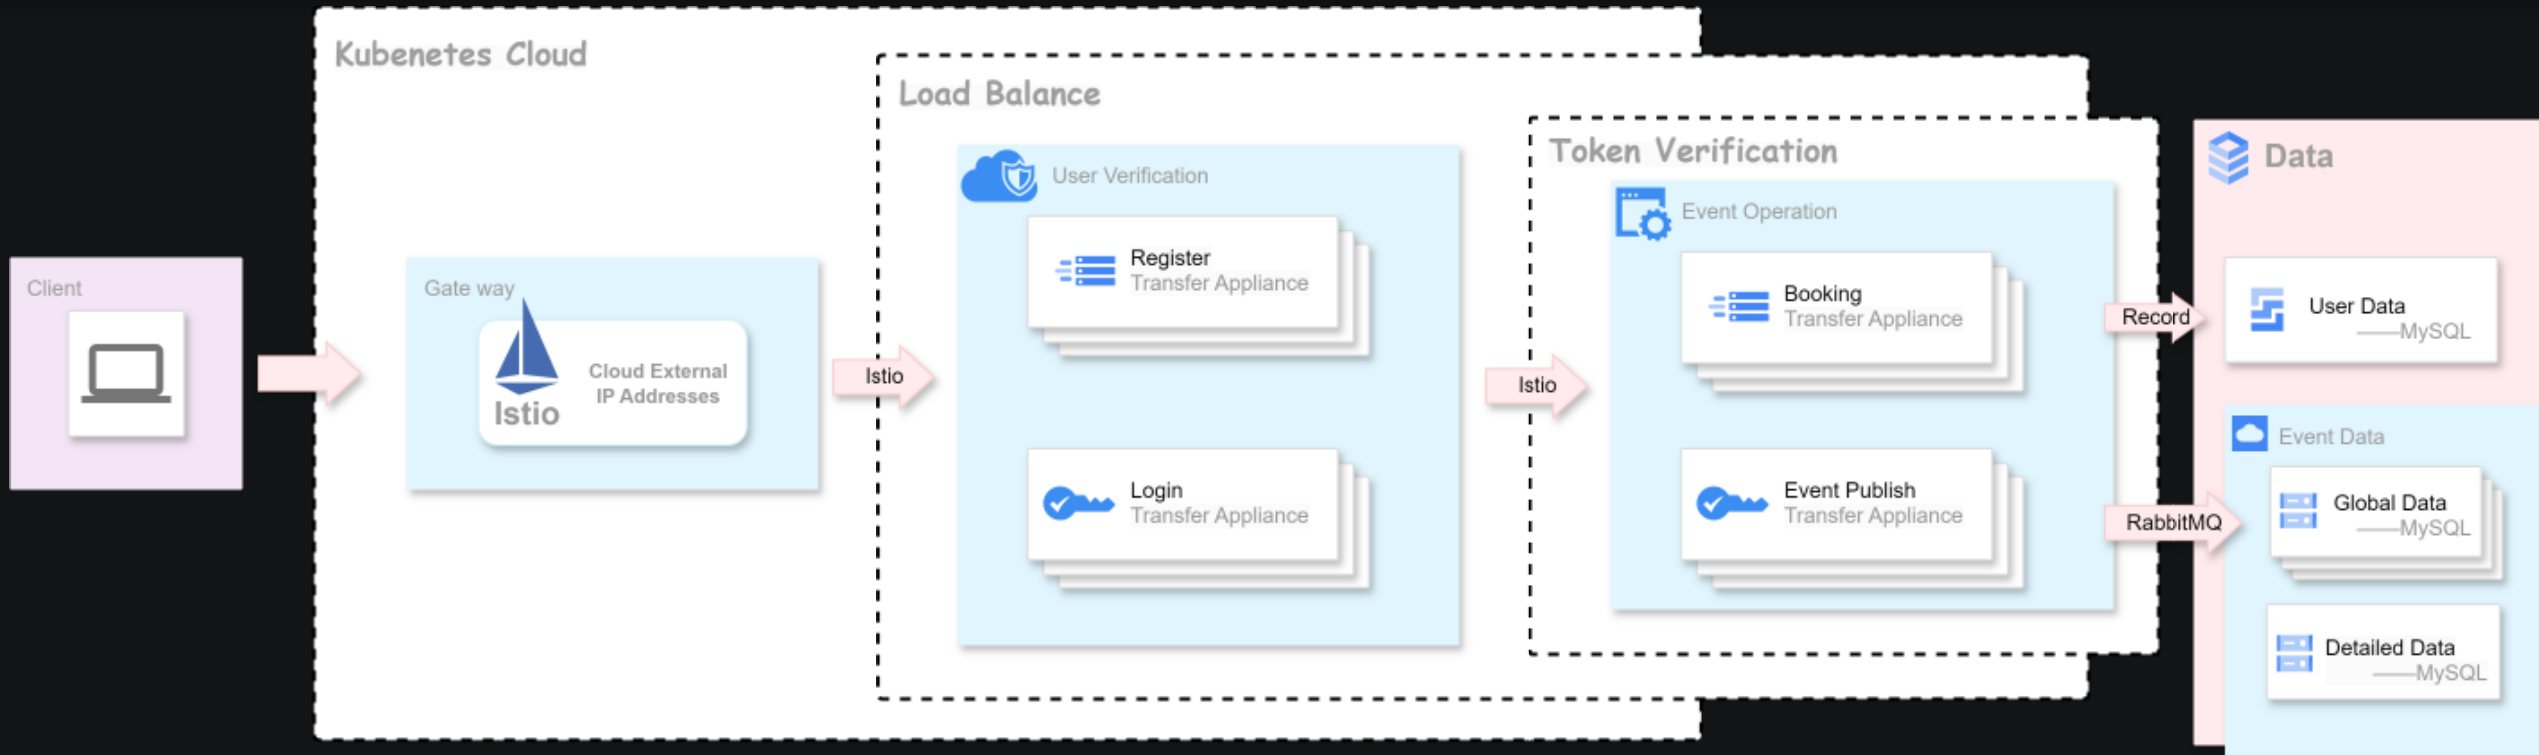
eventPublishService

Data Storage

eventGlobalDataService


eventDetailedDataService

System sturucture








Consistent Hashing

- We use **consistent hashing** .
 - **Why** we considered it?
- For the reason : **load balancing problem**



Consistent Hashing Principle

- e.g. 3 Deployment.



- entire hash value space into a virtual ring

- Next, 3 server deployment to the corresponding positions on the hash ring

- Well-integrated with Kubernetes

- Core Feature: **Server Scaling**



Service mesh and Istio

Service mesh

- A dedicated infrastructure layer to handle **service-to-service** communication in a microservices architecture
- Secure, fast, and reliable
- A set of lightweight network proxies
- Sidecar pattern

Istio



One of the most popular service mesh implementations



Ingress vs. Istio

Ingress

Purpose: Focus on routing from external networks

Implementation: Requires an Ingress controller (e.g., Nginx)

Configuration: Configured through Kubernetes resource objects (Ingress)

Istio

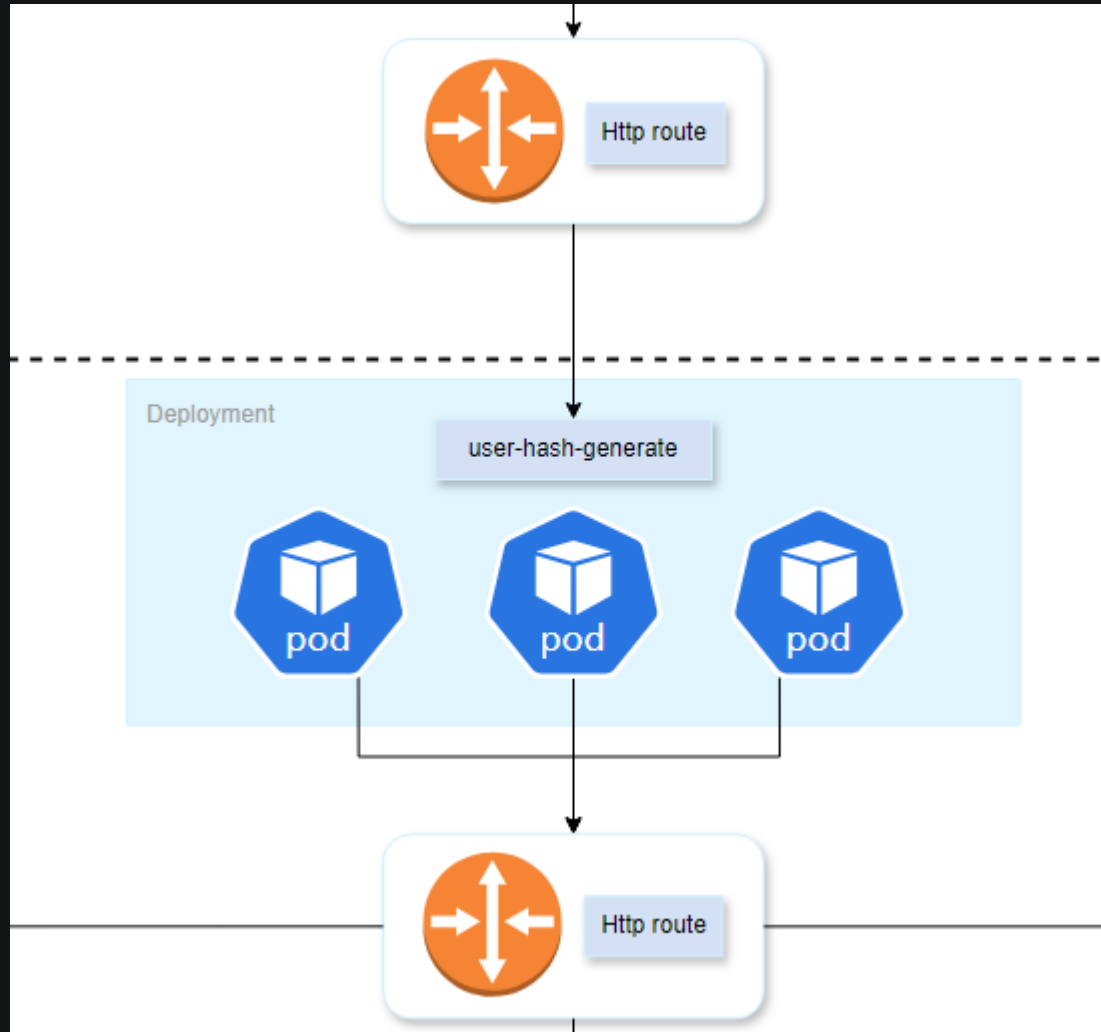
Purpose: A comprehensive service mesh solution

Implementation: Injecting a Sidecar proxy (Envoy Proxy) into Pods

Configuration: Through a series of CRDs



Hash generating and routing



Load balancer

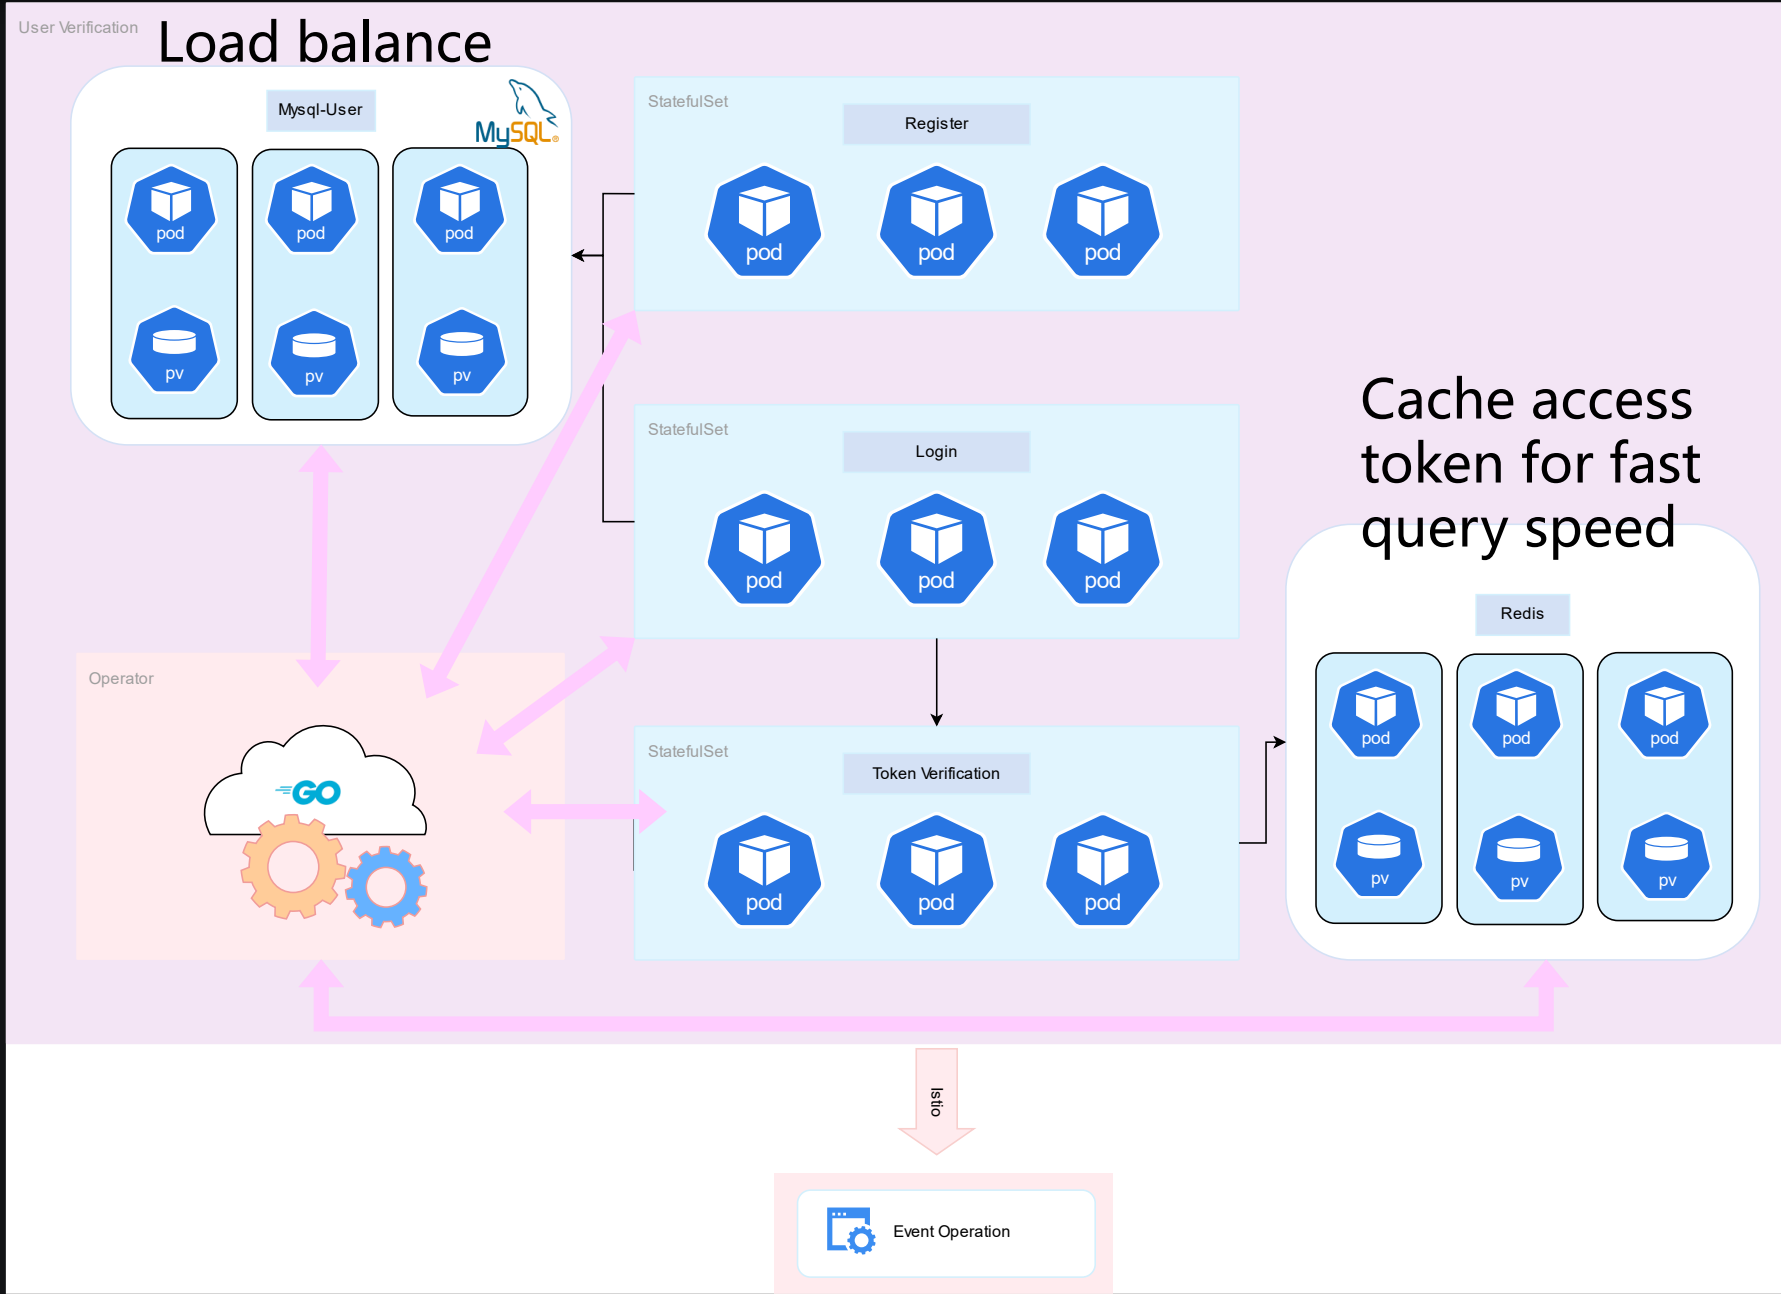
Generate hash based on the email address of users

Every email address will be mapped to a hash of length 8

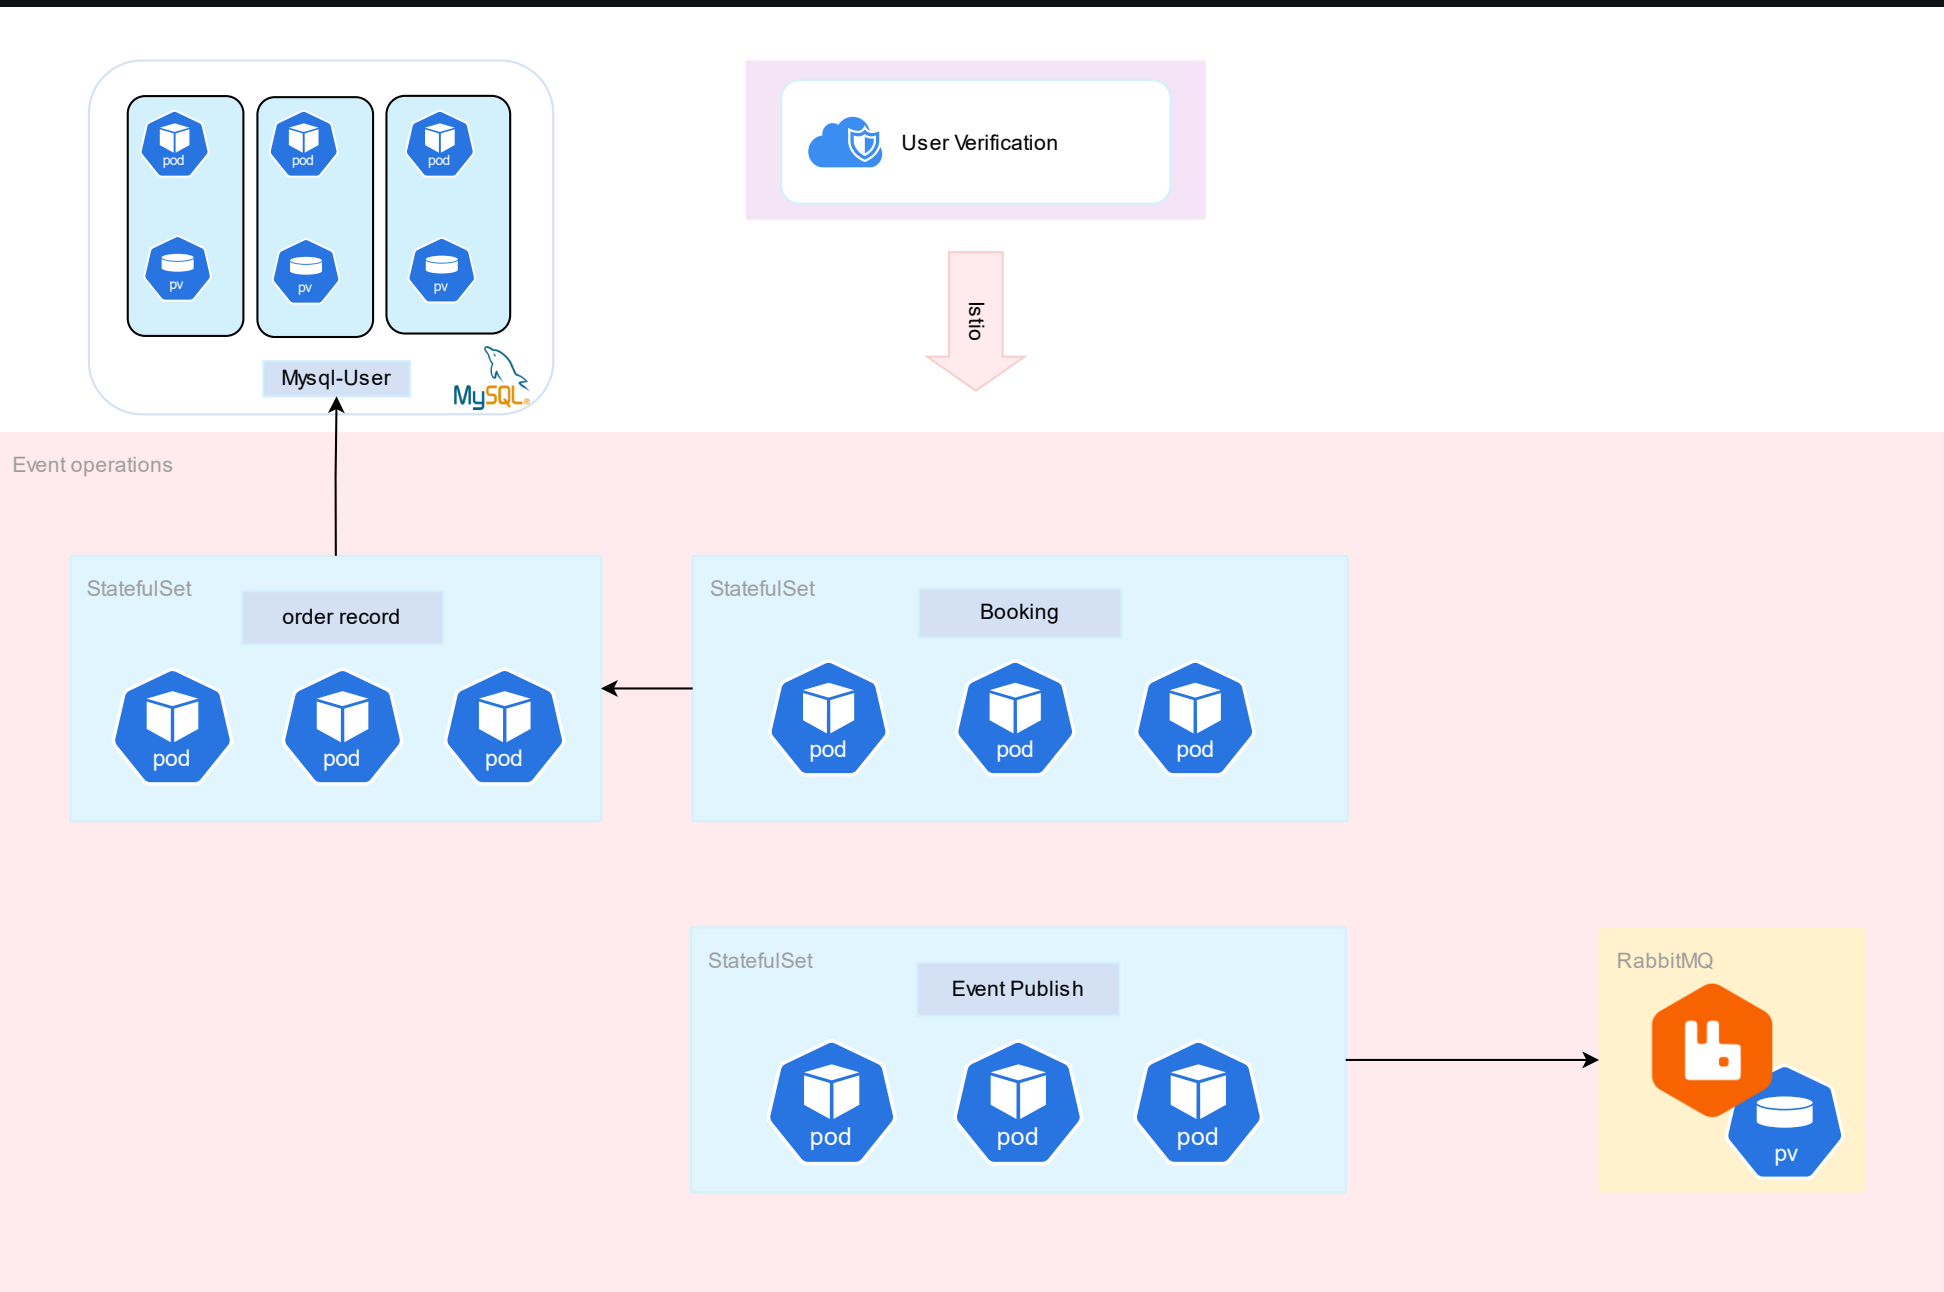
Routing based on hash



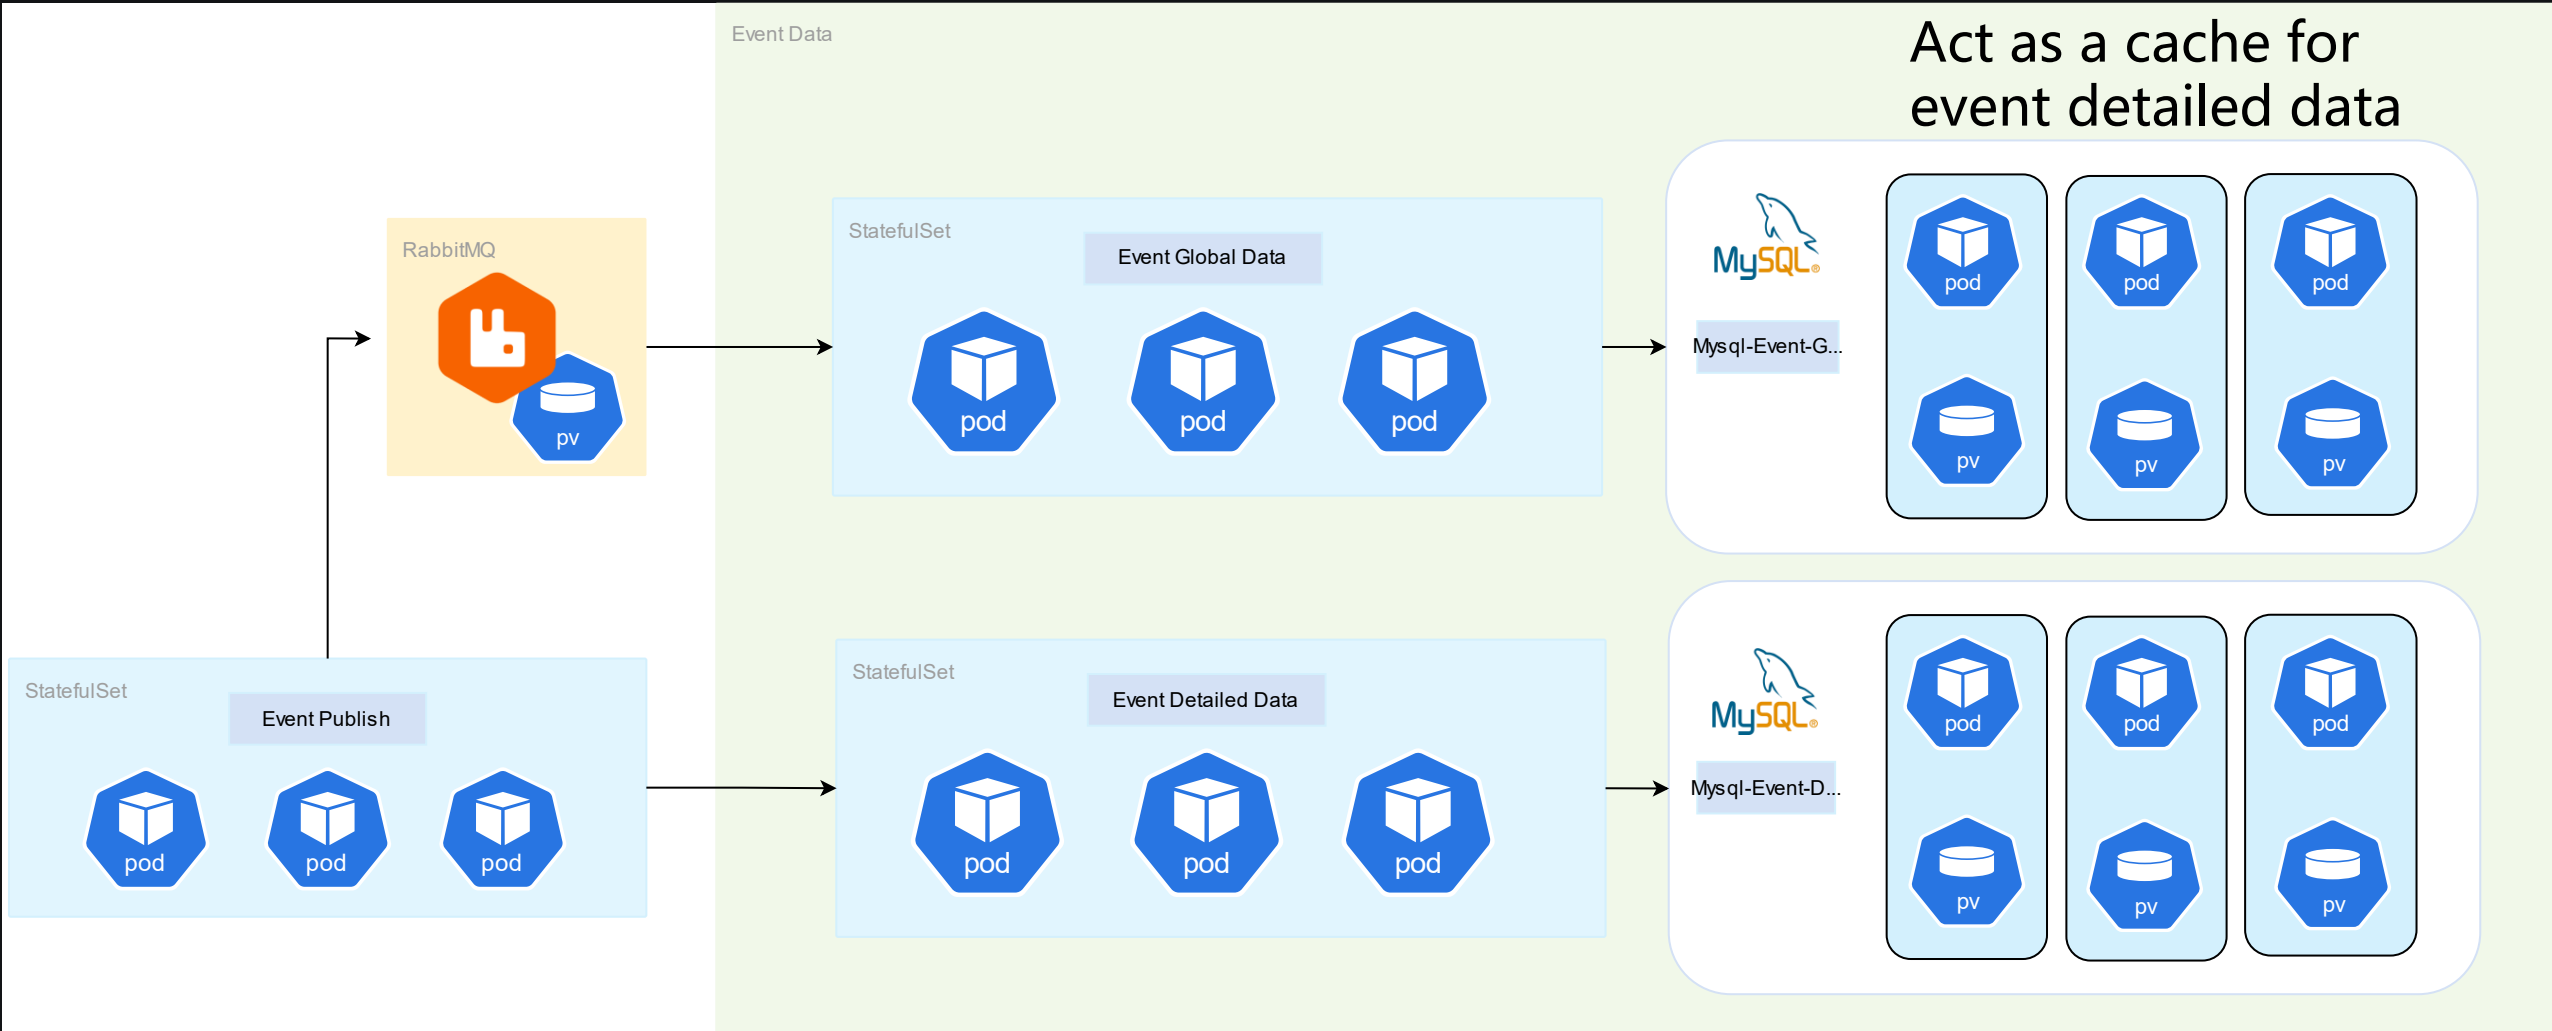
User Verification



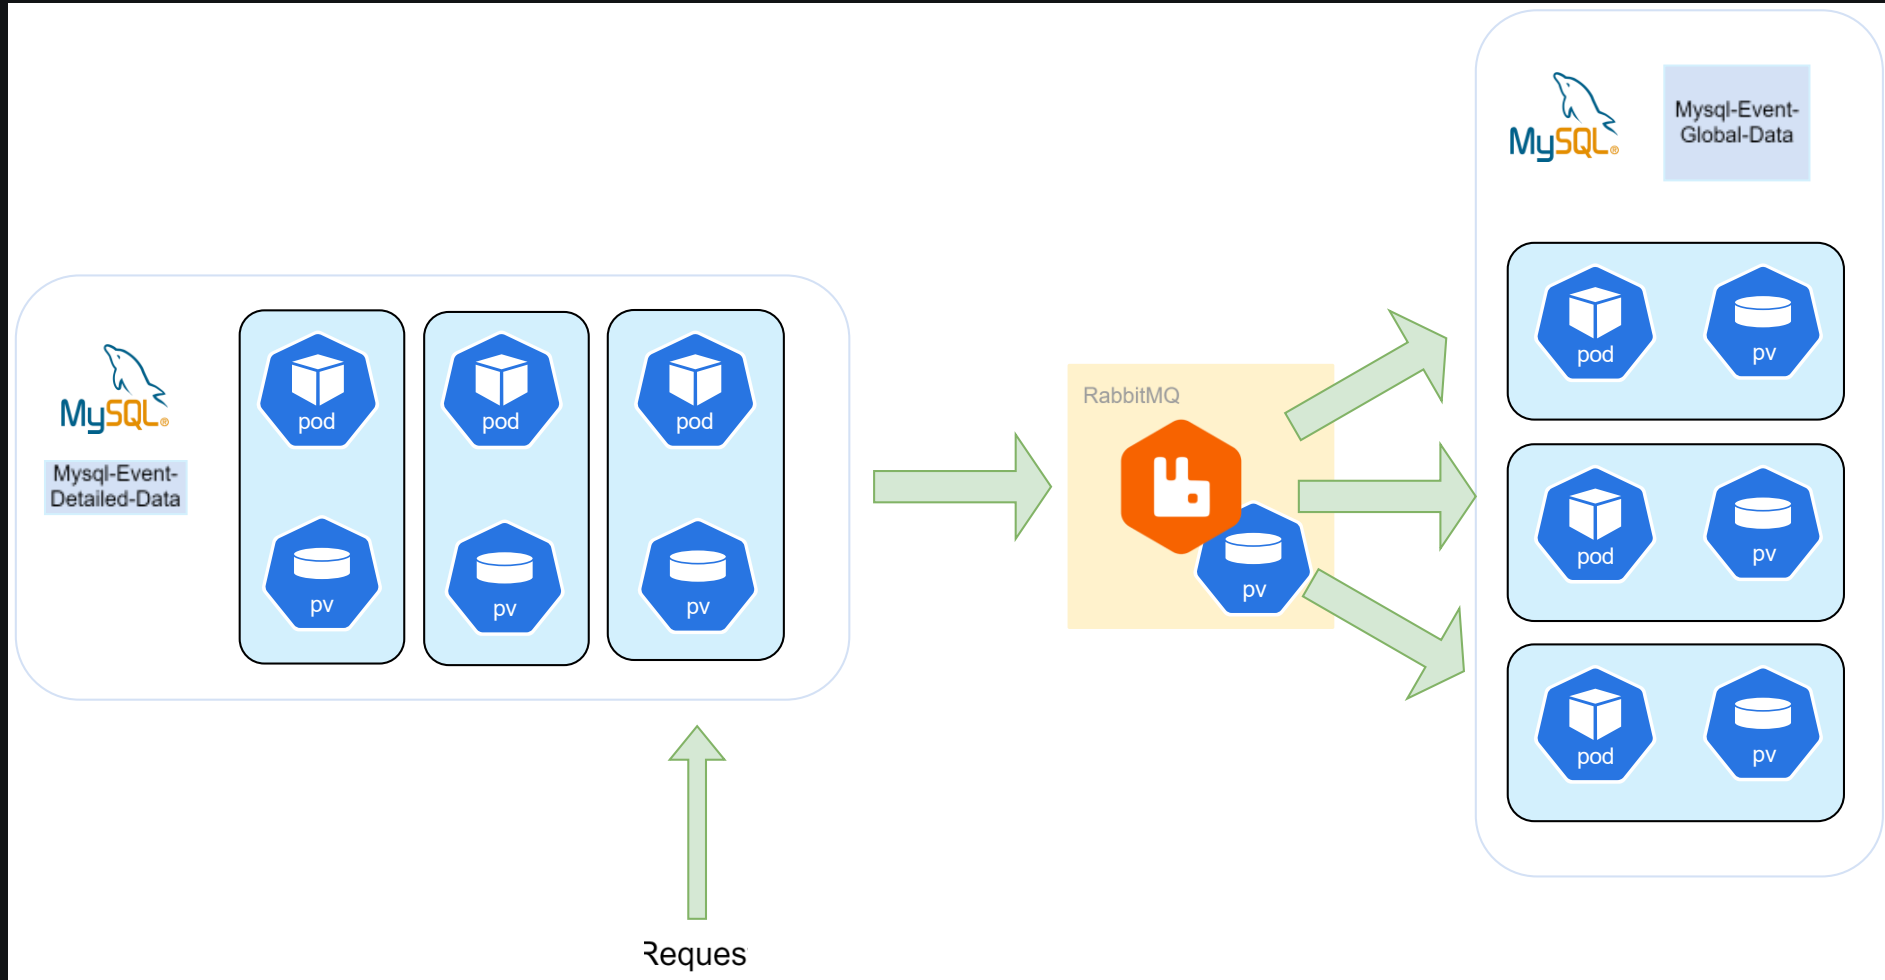
Event Operations



Data Storage



Data Synchronization




An Operator Implement by Our Own

Main purpose: Enable auto-scaling of our distributed databases

```
apiVersion: "auto-scaler.nus-cloud-project/v1"
kind: AutoScaler
metadata:
  namespace: nus-cloud-project
  name: test-auto-scaler
spec:
  maxTableSize: 10
```

A CRD to config the operator



Watching data in databases, split it if the table size exceeds the limit


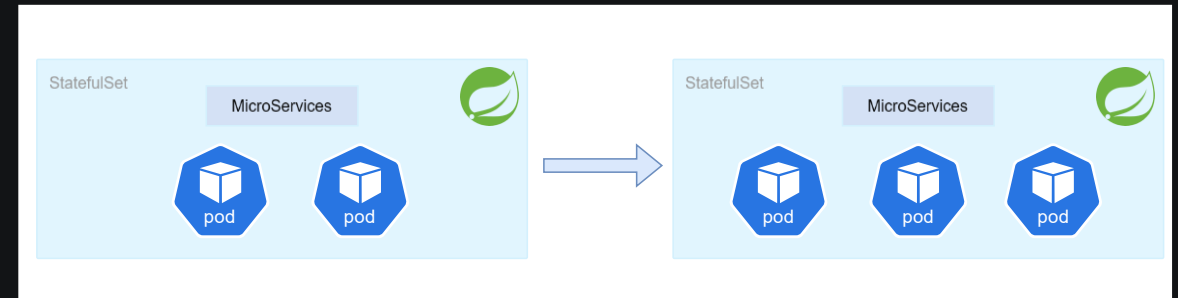

An Operator Implement by Our Own

1. Create new databases

2. Synchronize data to new databases

3. Create new microservices

4. Update routing rules



Our Workload


There are:

9 microservices

57 Kubernetes YAMLs

1 operator implemented by our own
with about **1000** lines of Go code


Our Code



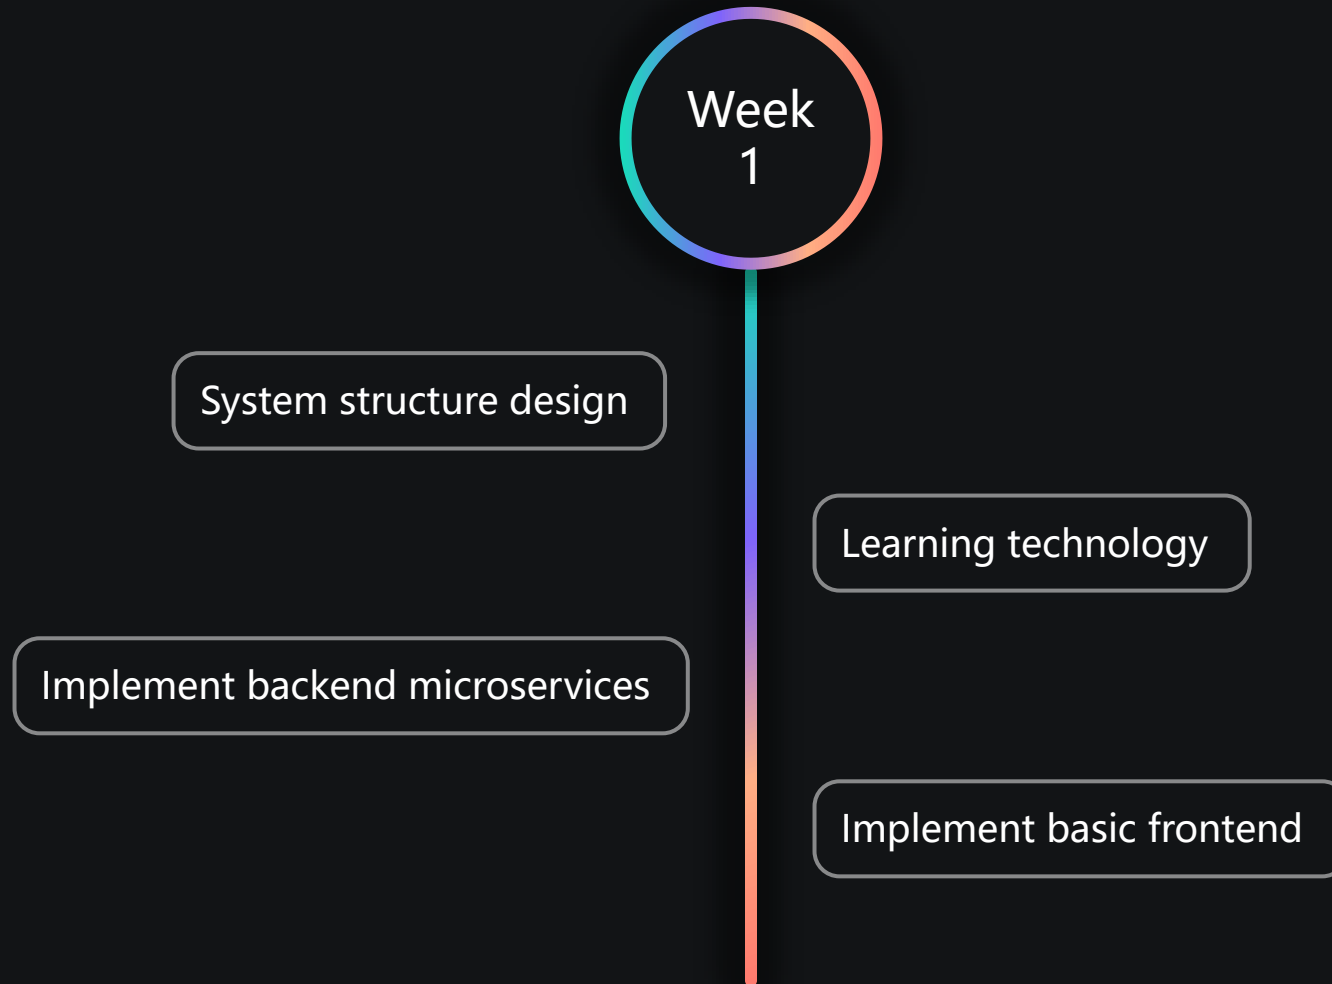
Demonstration Time



Poster



Project Timeline



Preliminary Work

Week
2

Implement system
structure on k8s

Implement
routing by hash

Implement
routing using Istio

Implement data
synchronization
using RabbitMQ

Week
3

Middle Presentation

Main Development

Week
3

Implement auto scaling
using operator

Improve frontend

Poster and ppt

Final Presentation



Any Questions?

Yes

No!



Thanks